

Stale Queues, Noisy Sizes: Can LLMs Discover Better Dispatch Strategies?

Kai Xu

New York University

Abstract

Real-world load-balancing problems often diverge from textbook setups: queue-length reports may be stale, job-size estimates may be unreliable, or a standard algorithm may degrade under a problem-specific failure mode that requires a novel correction. When practitioners encounter such scenarios, they increasingly turn to large language models for help—yet it is unclear whether LLMs can reason correctly about the stochastic properties that make these problems hard. We investigate this through two controlled experiments in which six LLMs zero-shot write dispatch functions for a server farm under heartbeat staleness and under noisy SRPT routing respectively. Across both experiments: no model reliably beats the best classical baseline in median performance, and naive deployment of LLM-generated functions frequently performs *worse* than the algorithm being replaced; scale determines code validity below roughly 8B parameters but does not predict median quality above that threshold; larger models occasionally discover high-quality solutions that smaller models never reach, even when their typical output is no better; and providing an explicit description of the failure mode increases the rate of correct solutions but still does not produce reliable improvements. These findings suggest that LLMs carry useful latent knowledge about algorithmic design but cannot yet be trusted to improve a deployed load-balancing algorithm without human oversight.

1 Introduction

A dispatcher routing jobs to a server farm always operates on imperfect information. Two failure modes motivate this paper. First, queue-length reports may be stale: by the time a heartbeat arrives from a worker, that worker has already processed additional jobs, and routing greedily on the stale report overloads whichever worker happened to report a short queue. Second, even with fresh state, the dispatcher may not know how long an incoming job will take and must rely on a noisy size estimate; the standard SRPT heuristic degrades when estimates are severely wrong.

Both scenarios reflect a common practical situation: an engineer knows that a standard algorithm applies but suspects it is underperforming, and turns to an LLM for a better alternative. Recent work shows LLMs can discover non-obvious algorithmic strategies when suitably prompted [Fawzi et al., 2024, Novikov et al., 2025, Ma et al., 2024], but these successes typically involve iterative search or explicit feedback. We ask whether zero-shot prompting alone is sufficient, and whether the answer changes when the LLM is told what is breaking. We test this with six models across two experiments: Experiment 1 on age-aware dispatch under heartbeat staleness, and Experiment 2 on adapting SRPT routing when job-size estimates are heavily noisy.

Our main findings are:

1. **No model reliably beats the best classical baseline.** In Experiment 1, every model’s median performance is worse than routing uniformly at random—the weakest possible baseline. In Experiment 2, median performance sits between Least Connections and SRPT-noisy;

only a small minority of trials improve on SRPT-noisy. Naively deploying LLM-generated functions can produce results substantially worse than the algorithm being replaced.

2. **Scale gates code validity but not median quality.** Below roughly 8B parameters, models rarely produce valid, executable dispatch functions. Above that threshold, larger models do not achieve better median performance: the largest Qwen3 model (14B) does not outperform smaller models in either experiment.
3. **Scale does affect the extremes.** Larger models occasionally discover high-quality solutions that smaller models never find, even when their typical output is no better. The correct strategy in each experiment appears in only a small fraction of trials, concentrated in the larger models.
4. **Describing the failure mode helps, but is not sufficient.** Experiment 2’s prompt explicitly names the phantom-underload failure; the rate of correct solutions is much higher than in Experiment 1, but median performance still does not reliably exceed the best baseline.

2 Related Work

Load balancing under staleness. Join-Shortest-Queue (JSQ) [Lin and Raghavendra, 1996] is optimal under complete, instantaneous state information. Under stale reports, JSQ is susceptible to the thundering herd: all dispatchers simultaneously route to whoever last reported a short queue, overloading that worker until the next heartbeat resets the picture. Power-of-Two-Choices [Mitzenmacher, 2001] disrupts herding by sampling two workers and routing to the shorter, trading information for robustness. Mitzenmacher [2000] shows that even moderately outdated information can eliminate Po2C’s guarantees. Dahlin [2000] characterizes the herding failure analytically and derives age-adjusted routing thresholds—the same kind of correction our experiments ask LLMs to discover independently.

SRPT scheduling and estimation noise. Schrage [1968] proved that preemptive SRPT minimizes mean response time in the M/G/1 queue; this optimality requires exact job sizes. In practice, job sizes must be estimated, and estimation errors degrade performance. Heavy-tailed service times, common in real workloads [Crovella and Bestavros, 1997], amplify the damage because large jobs are common enough to be persistently present, and a severe underestimate on a large job creates a long-lasting phantom underload.

LLMs for algorithm synthesis. Fawzi et al. [2024] and Novikov et al. [2025] use LLM-driven search to discover improvements to combinatorial optimization and data-center scheduling, achieving substantial gains over human-designed baselines. Ma et al. [2024] synthesizes reward functions for reinforcement learning zero-shot from problem descriptions alone. Cheng et al. [2025a] examine LLM-generated heuristics for scheduling problems in a systems context. Our work differs in two respects: correctness is determined by steady-state simulation rather than test cases or human evaluation, and we study zero-shot synthesis without iterative search or feedback, which represents the more common case when a practitioner queries an LLM for algorithm advice.

AI for systems. Sparrow [Ousterhout et al., 2013] implements Po2C in a real cluster scheduler, showing the two-choices principle works at production scale. Cheng et al. [2025b] surveys AI-driven systems research and reports $5\times$ speedups on load-balancing tasks via LLM-guided search. Our

setting is simpler than production cluster scheduling, which lets us isolate the reasoning challenge cleanly: one dispatch decision per job, one integer output, full-simulation evaluation.

3 Experiment 1: Age-Aware Dispatching under Heartbeat Staleness

3.1 System Model

We simulate a single dispatcher routing to $N=8$ workers. Jobs arrive as a Poisson process at total rate $\lambda=6.4$ jobs/ms ($\rho=0.8$). Service times are i.i.d. Exponential(mean 1 ms). Each worker maintains an independent FIFO queue.

Each worker emits a heartbeat every $H=20$ ms. Each heartbeat incurs a propagation delay drawn from Uniform(1, 5) ms. The dispatcher records the reported queue length and the arrival time of the most recent heartbeat. At dispatch time t , worker w 's heartbeat age is $a_w = t - \text{last_heard}_w$, ranging from near zero to nearly $H+5$ ms. Mean age at dispatch is approximately 11–12 ms, or roughly 11–12 expected completions of staleness.

```
def dispatch(queue_lengths: list[int],
            queue_ages_ms: list[float],
            num_workers: int) -> int
```

`queue_lengths[w]` is the stale reported queue depth; `queue_ages_ms[w]` is the age in milliseconds of the most recent report. The function returns the worker index to receive the incoming job.

Metric. We sample the true queue depths every 10 ms and compute the cross-worker standard deviation σ . We report mean σ over 10^5 jobs; lower is better. Each strategy is evaluated over 100 independent seeds.

3.2 Baselines

Three classical policies serve as reference points:

- **Random.** Route uniformly at random. Any strategy worse than Random is actively harmed by the information it uses.
- **Po2C (stale).** Sample two workers and route to the shorter reported queue [Mitzenmacher, 2001]. Disrupts herding at the cost of ignoring ages entirely.
- **JSQ (stale).** Route to the minimum reported queue, breaking ties randomly. Uses stale queue lengths greedily; highly susceptible to the thundering herd.

JSQ (stale) achieves median $\sigma=14.95$, nearly four times worse than Random ($\sigma=3.77$). Figure 1 shows the mechanism directly: one or two workers repeatedly spike above 50 jobs while others idle, resetting each heartbeat cycle. Po2C reduces imbalance to $\sigma=4.40$ through randomization but leaves age information entirely unused.




Figure 1: Thundering herd under JSQ-stale. Per-worker true queue depth over 400 ms (seed 0). One or two workers repeatedly spike above 50 jobs while the rest idle. Each spike resets when a heartbeat reassigns the apparent shortest-queue label. JSQ-stale ($\sigma=14.95$) performs nearly four times worse than Random ($\sigma=3.77$).

3.3 LLM Evaluation

We evaluate six models: Qwen3-1.7B, Qwen3-4B, Qwen3-8B, Qwen3-14B, DeepSeek-R1-0528-Qwen3-8B (hereafter DeepSeek-R1), and gpt-oss-20b. Each model is prompted at temperature 1 to write the body of the `dispatch` function, given the system parameters and interface docstring. No strategy is suggested, no examples are provided, and the thundering herd is not named. Responses are parsed by extracting content between `<code>` tags.

Each generated function is paired with a fixed simulation seed and evaluated for 100 independent trials per model. A trial is valid if the code parses, executes without error on every dispatch call, and returns an integer in $[0, N)$. Parse rates range from 20% to 84%, yielding 257 valid trials in total (Table 1).

3.4 Results

No model reliably beats Random. Table 1 shows median mean σ across valid trials. Every model performs substantially worse than Random. The best median is 21.0 (gpt-oss-20b), $5.6\times$ worse than Random and $4.8\times$ worse than Po2C. There is no clear scaling relationship: Qwen3-8B (24.0) outperforms the larger Qwen3-14B (27.4), and gpt-oss-20b achieves the best median despite the lowest parse rate. Figure 2 shows the per-trial distribution for each model.

Table 1: Experiment 1 results. Median mean σ over 100 seeds for baselines and over all valid trials for LLMs. Lower is better. [†]Reported parameter count from model metadata.

Strategy / Model	Params	Parse%	n	Med. σ
<i>Classical baselines (100 seeds)</i>				
Random	—	—	100	3.77
Po2C (stale)	—	—	100	4.40
JSQ (stale)	—	—	100	14.95
<i>LLM-generated dispatchers</i>				
Qwen3-1.7B	1.8B	84%	82	27.21
Qwen3-4B	4.0B	55%	40	26.75
Qwen3-8B	8.5B	43%	36	24.02
Qwen3-14B	15.7B	68%	60	27.38
DeepSeek-R1 (8B)	8.5B	45%	29	27.26
gpt-oss-20b	20B [†]	20%	10	21.03




Figure 2: Per-trial σ distributions in Experiment 1. Each dot is one valid trial; the red bar marks the within-model median. Most models cluster in the $\sigma \approx 20$ – 35 range. gpt-oss-20b shows a bimodal distribution: a small cluster of trials near the Random reference line alongside the majority well above it, corresponding to the rare trials that implement age subtraction correctly.

The age-aware subtraction strategy was independently discovered. Four trials (three from gpt-oss-20b, one from Qwen3-14B) achieve $\sigma < 3.77$, beating Random. All four implement the

same strategy:

$$\text{adj}_w = \max(q_w - a_w, 0), \quad \text{route to } \arg \min_w \text{adj}_w \quad (1)$$

Since the expected number of completions at worker w in a_w ms is $a_w/\mu = a_w$ (with $\mu=1$), subtracting a_w from the stale queue length estimates the current queue depth. Clamping at zero handles negative estimates. These four trials achieve $\sigma \approx 2.9$, below Random. Figure 3 shows one representative from each model.

gpt-oss-20b, trial 53
mean $\sigma = 2.87$

Qwen3-14B, trial 84
mean $\sigma = 2.98$

```

min_index = 0
min_value = float('inf')
for w in range(num_workers):
    eff = queue_lengths[w] - queue_ages_ms[w] / 1.0
    if eff < 0:
        eff = 0
    # Add small random noise to break ties uniformly
    eff += random.random() * 0.01
    if eff < min_value:
        min_value = eff
        min_index = w
return min_index

```

```

adjusted = []
for w in range(num_workers):
    adj_len = max(queue_lengths[w] - queue_ages_ms[w], 0)
    adjusted.append((adj_len, w))
min_adj = min(adjusted, key=lambda x: x[0])
candidates = [w for (len_, w) in adjusted if len_ == min_adj[0]]
return random.choice(candidates)

```

Figure 3: Independent convergence on the age-subtraction strategy. Two models from different families produce the same dispatch logic without coordination: compute $\text{adj}_w = \max(q_w - a_w, 0)$, return $\arg \min_w \text{adj}_w$. Both achieve $\sigma \approx 2.9$, below Random’s 3.77.

Strategy classification. We classify each valid trial by inspecting the generated code (Figure 4):

- **Subtract age** ($n=12$, median $\sigma=20.9$): implements $q_w - f(a_w)$ for some positive f . All four trials that beat Random belong to this class.
- **JSQ + age tiebreaker** ($n=96$, median $\sigma=26.7$): routes by shortest queue; age breaks ties only. Effectively reduces to JSQ-stale in practice.
- **Add age** ($n=49$, median $\sigma=22.1$): computes $q_w + \alpha \cdot a_w$. Interprets a long silence as evidence of load rather than as staleness, reversing the correct direction.
- **Other** ($n=100$, median $\sigma=32.7$): miscellaneous heuristics with no consistent pattern.

All four classes perform worse than Random. The key observation is not which direction models get the sign wrong, but that *any* use of the age signal—in either direction—consistently fails to beat the baseline that ignores it entirely. Using stale information badly is worse than not using it at all.




Figure 4: Performance by strategy class in Experiment 1. Each dot is one valid trial; vertical bars show the within-class median. The subtract-age class contains the four trials that beat Random. Dashed lines mark the Random and JSQ-stale medians.

4 Experiment 2: Noisy SRPT Dispatching with Heavy-Tailed Jobs

4.1 System Model

The second experiment replaces heartbeat staleness with a different source of uncertainty: noisy job-size estimates. A single dispatcher routes jobs to $N=8$ workers at utilization $\rho=0.8$ ($\lambda=6.4$ jobs/ms). Service times follow a Pareto distribution with shape $\alpha=1.5$ and mean 1 ms. The Pareto distribution is heavy-tailed: the probability of a job exceeding x times the mean decays as $x^{-\alpha}$, so large jobs are substantially more common than under exponential service and concentrate a disproportionate share of total work.

Each incoming job comes with a noisy size estimate $\hat{S} = S \cdot \epsilon$, where S is the true size and $\epsilon = \exp(Z)$, $Z \sim \mathcal{N}(0, \sigma^2)$, $\sigma=1.0$. This multiplicative noise has geometric standard deviation $e^\sigma \approx 2.7$: estimates within one log-scale standard deviation range from roughly $e^{-1} \approx 0.37\times$ to $e^{+1} \approx 2.7\times$ the true size.

The dispatcher maintains, per worker, an accumulated estimated remaining work (ERW):

$$\text{ERW}_w(t) = \sum_{j \in Q_w(t)} \hat{S}_j \quad (2)$$

where $Q_w(t)$ is the set of jobs currently at worker w . This sum is updated when jobs arrive and depart but does not decrease during service; the estimate is frozen at arrival.

```
def dispatch(connections: list[int],
             estimated_remaining_work: list[float],
             job_size_estimate: float,
             num_workers: int) -> int
```

`connections[w]` is the exact active job count at worker w (no noise, no delay). `estimated_remaining_work[w]` is ERW_w , a noisy accumulated estimate. The function returns the worker index to receive the incoming job.

Metric. We measure mean response time in milliseconds across 10^5 jobs, averaged over seeds. Lower is better. A well-designed dispatcher should achieve lower mean response time than naive SRPT with noisy estimates, ideally approaching the oracle.

4.2 How Noisy Estimates Break SRPT Routing

SRPT is optimal for a single M/G/1 queue [Schrage, 1968], but the optimality requires exact job sizes. In our N -worker setting, routing each arrival to the worker with the smallest estimated total remaining work approximates SRPT; under noisy estimates this approximation fails in two ways. First, because ϵ is log-normal, individual estimates can be far from the true size: at $\sigma=1$, roughly 1% of estimates fall below 10% of the true value. Second, ERW_w does not update during service—a large underestimated job contributes only its small estimated size to ERW_w throughout its entire service time.

Concretely, suppose job j has true size $S_j=50$ ms but estimate $\hat{S}_j=0.5$ ms. For the 50 ms duration of j 's service, ERW_w underreports worker w 's load by approximately 49.5 ms. The dispatcher routes all arrivals to w because it appears nearly idle. We call this a *phantom underload*. The Pareto distribution ($\alpha=1.5$) amplifies the damage: jobs larger than $5\times$ the mean contribute approximately 26% of total work, so phantom underloads are both frequent and long-lasting.

Figure 5 shows the degradation empirically. SRPT-noisy (using noisy estimates for routing) degrades smoothly with σ from near-oracle performance at $\sigma=0$ to 10% above oracle at $\sigma=2$. Least Connections (LC), which ignores size estimates entirely and routes to the worker with the fewest active jobs, is unaffected by σ and serves as a noise-robust reference. Random is likewise σ -independent, as expected.




Figure 5: SRPT degrades with estimation noise. Median mean response time (50 seeds) versus log-normal noise σ for four baseline strategies. SRPT-oracle uses true job sizes and serves as a performance ceiling. SRPT-noisy degrades monotonically with σ . Least Connections and Random are unaffected because they do not use size estimates. The vertical line at $\sigma=1.0$ marks the operating point used in the LLM evaluation.

The ideal correction. The SRPT failure stems from ERW_w systematically underestimating the load on workers with underestimated large jobs. An alternative, noise-free signal for remaining work is $c_w \cdot E[S]$, where $c_w = \text{connections}[w]$ is the exact active job count and $E[S]=1$ ms is the known mean service time. This signal cannot be made arbitrarily small by underestimation because it depends only on counts, not on size estimates. A corrected score that blends both signals,

$$\tilde{s}_w = ERW_w + \hat{S}_{\text{new}} + c_w \cdot E[S], \quad (3)$$

adds a floor term proportional to the number of active jobs. When ERW_w is severely underestimated, the floor term prevents worker w from appearing idle, limiting the accumulation of phantom-underload traffic. An equivalent variant enforces a hard floor: $\tilde{s}_w = \max(ERW_w, c_w \cdot E[S]) + \hat{S}_{\text{new}}$, which discards ERW_w entirely when the connection-count estimate is larger.

4.3 Baselines

We compare four reference policies, each evaluated over 20 seeds:

- **SRPT-oracle.** Route to $\arg \min_w (\text{TRW}_w + S_{\text{new}})$, where TRW_w is the true total remaining work using exact job sizes. Achieves median 1.79 ms; serves as a performance ceiling.

- **SRPT-noisy.** Route to $\arg \min_w (\text{ERW}_w + \hat{S}_{\text{new}})$, using noisy estimates throughout. Achieves median 1.94 ms.
- **Least Connections (LC).** Route to $\arg \min_w c_w$, breaking ties randomly. Ignores size estimates entirely; achieves median 2.35 ms.
- **Random.** Route uniformly at random. Achieves median 35.8 ms; a weak baseline included to confirm that LLM-generated functions are not trivially degenerate.

4.4 LLM Evaluation

We use the same six models and the same prompting protocol as Experiment 1, adapted to the new interface. The prompt describes the SRPT failure mode explicitly: it explains that ERW_w is a frozen accumulated sum, not a running estimate, and that severely underestimated jobs make a worker appear nearly idle throughout their service. We do not suggest a correction strategy.

We generate 100 trials per model. A trial is valid if the function parses, executes without error, and does not fall back to a baseline dispatcher on every call (fallback rate < 1.0). We exclude fallback-dominated trials because their measured performance reflects the fallback strategy, not the LLM’s intended logic. Table 2 reports parse rates and valid trial counts after this filter.

Table 2: Experiment 2 results (valid trials after filtering fallback-dominated runs; models with $n < 5$ are reported for completeness but are too small for reliable inference). LLM rows are ordered by valid trial count. Baselines are evaluated over 20 independent seeds. Lower median mean response time is better.

Strategy / Model	Parse%	n valid	Med. RT (ms)	IQR (ms)
<i>Baselines (20 seeds each)</i>				
SRPT-oracle	—	—	1.79	[1.55, 2.10]
SRPT-noisy	—	—	1.94	[1.66, 2.26]
LC	—	—	2.35	[1.91, 2.66]
Random	—	—	35.8	[25.0, 76.4]
<i>LLM-generated dispatchers</i>				
Qwen3-1.7B	1%	0	—	—
Qwen3-8B	6%	1	1.65	—
Qwen3-4B	26%	6	2.47	[1.76, 3.16]
DeepSeek-R1 (8B)	27%	16	2.03	[1.88, 2.41]
gpt-oss-20b	36%	30	2.02	[1.76, 3.15]
Qwen3-14B	50%	50	2.44	[1.93, 3.49]

4.5 Results

Scale determines code validity, not performance rank. Parse rates and valid trial counts broadly increase with model scale, though not monotonically: Qwen3-1.7B produces no valid simulations. Qwen3-4B (26% parse rate, 6 valid trials) outperforms the larger Qwen3-8B (6%, 1 trial), suggesting that architecture and training matter alongside raw parameter count. gpt-oss-20b and DeepSeek-R1 reach 27–36% parse rates with 16–30 valid trials, and Qwen3-14B achieves the highest parse rate at 50% (50 valid trials), giving sufficient sample sizes for inference.

Among models with adequate samples, the ranking does not follow parameter count. gpt-oss-20b (median 2.02 ms) and DeepSeek-R1 (2.03 ms) achieve similar performance close to SRPT-noisy

(1.94 ms), while the larger Qwen3-14B (2.44 ms) sits closer to LC. Figure 6 shows the full per-model distributions.




Figure 6: Per-model response time distributions in Experiment 2. Box plots show the spread of mean response time across valid trials (100 seeds each, fallback-dominated trials excluded). Dashed reference lines mark the four baseline medians. gpt-oss-20b and DeepSeek-R1 cluster near SRPT-noisy; Qwen3-14B shows wider variance. Models with $n < 5$ valid trials are omitted.

Models converge on two distinct strategies. Every valid trial falls into one of two strategy classes. *Both-signals* functions (84 trials, median 2.06 ms) use a combination of ERW_w and c_w , corresponding to some form of the correction in Equation (3). *LC-only* functions (19 trials, median 2.58 ms) route solely by connection count, effectively reimplementing the LC baseline. No valid trial implements a pure SRPT-only strategy, despite the prompt showing SRPT as the starting point. Models told that SRPT’s estimates are unreliable consistently chose to down-weight or abandon the size-estimate signal; the better outcome is to correct it rather than discard it, which the both-signals class does.

The best functions independently discover bias correction. Sorting all 103 valid, non-fallback trials by mean response time and removing exact source-code duplicates yields five unique top-performing functions, all belonging to the both-signals class. Two structurally distinct variants appear repeatedly:

- **Additive penalty** (ranks 1, 3, 4; primarily from gpt-oss-20b): $\tilde{s}_w = \text{ERW}_w + \hat{S}_{\text{new}} + c_w \cdot 1.0$
- **Floor correction** (ranks 2, 5; primarily from Qwen3-14B): $\tilde{s}_w = \max(\text{ERW}_w, c_w \cdot 1.0) + \hat{S}_{\text{new}}$

Both use the coefficient 1.0 ms, equal to $E[S]$ —the mean service time stated in the prompt. The additive penalty blends both signals; the floor correction replaces ERW_w with $c_w \cdot E[S]$ whenever the latter is larger. These functions independently reconstruct the correction derived in Section 4.2 from first principles.

To evaluate robustly, we re-run each of the five unique functions over 20 fresh seeds (the same 20 seeds used for the baselines in Table 2). The top two functions achieve median mean response times of 1.88 ms and 1.90 ms respectively—between SRPT-oracle (1.79 ms) and SRPT-noisy (1.94 ms). Figure 7 shows the code alongside the performance comparison.




Figure 7: Best LLM-discovered dispatch functions. *Left:* Source code for the two top-performing unique functions. Both independently discover a connection-count correction to the SRPT score; they differ only in whether the correction is additive or acts as a floor. *Right:* Robust performance comparison (20 seeds each) of all five top functions against the four baselines. The best two functions fall between SRPT-oracle and SRPT-noisy. Error bars show the interquartile range.

These results are best-case: the functions shown are the top entries from a sorted list of 73 unique valid functions, and they were discovered inconsistently. gpt-oss-20b produced 30 valid trials; 3 of them rank in the top 5 while others perform near the LC baseline. Qwen3-14B produced 50 valid trials with wider variance. The typical performance of both models is 2.0–2.4 ms, not 1.88 ms. LLMs can synthesize the correct correction when given a description of the failure mode, but not reliably—the best functions were discovered in 3 of 30 trials for the top model, not consistently.

5 Discussion

No model reliably beats the best baseline. The central finding across both experiments is that LLM-generated dispatch functions do not reliably improve on the best available classical baseline. In Experiment 1, every model’s median performance is worse than Random—meaning the typical LLM output is actively harmed by the information it uses. In Experiment 2, median performance sits between LC and SRPT-noisy; improvements over SRPT-noisy are confined to the top few functions out of over a hundred trials. A practitioner who deployed a randomly sampled LLM-generated function would, in expectation, get worse performance than if they had kept the original algorithm.

The two experiments differ in *how bad* the typical output is. In Experiment 2 the task is more forgiving: routing by connection count alone is reasonable, so abandoning ERW is not catastrophic. In Experiment 1, using heartbeat age incorrectly is strictly worse than ignoring it, so more of the strategy space leads to actively harmful results. But in both cases, the median LLM output falls short of what a careful algorithm designer would produce.

Scale as a threshold, not a gradient. In both experiments, scale acts as a threshold rather than a continuous performance gradient. Models below roughly 8–14B parameters rarely produce valid, non-fallback dispatchers at all. Above that threshold, model size does not appear to correlate with median response time: Qwen3-14B, the largest Qwen3 model tested, does not outperform gpt-oss-20b in median performance despite having more parameters, and DeepSeek-R1 matches gpt-oss-20b without clearly exceeding it.

Where scale does seem to matter is at the extremes. Smaller models tend to generate more syntactically broken or semantically incoherent functions, pulling down their worst-case performance. Larger models occasionally discover high-quality solutions—diamonds in the rough—that smaller models never reach, even though their typical output is no better. This asymmetry suggests that scale increases the ceiling on best-case solutions without consistently raising the floor.

This pattern raises a natural question about the underlying cause. The inconsistency between occasional correct solutions and poor median performance may reflect a generalization failure in the reinforcement learning stage of these models’ training: RLHF or RLVR reward signals optimized on common code and reasoning tasks may not transfer reliably to out-of-distribution problems like stochastic load balancing. If so, the largest models carry latent knowledge of the correct strategy—visible in the rare trials that get it right—but cannot sample it consistently because the policy was not shaped by feedback from this problem class.

Describing the failure mode helps. The two experiments differ in one key respect beyond the failure mode itself: Experiment 2’s prompt explicitly names the phantom-underload mechanism, while Experiment 1’s prompt provides only the interface and system parameters. The best-discovered functions in Experiment 2 correct exactly the described failure, using a coefficient (1.0ms) equal to the stated mean service time—evidence that the model is reading and applying the failure description, not just pattern-matching to the interface. This asymmetry explains the gap in discovery rates: Experiment 2 yields more top-performing solutions despite presenting a harder optimization landscape. A clean ablation—add an explicit thundering-herd description to Experiment 1, or withhold the ERW-freeze explanation from Experiment 2—would isolate how much of the gap is attributable to prompt framing versus task difficulty.

Limitations. Valid trial counts for small models (Qwen3-4B: $n=6$; Qwen3-8B: $n=1$) are too small for reliable inference in Experiment 2 and should be treated as qualitative. All results are

from a single prompt template; different prompts may yield different strategy distributions. The simulation uses a fixed set of parameters ($N=8$, $\rho=0.8$, $\alpha=1.5$); generalization to other operating points is not tested.

Directions for future work. Several open questions follow directly from these results. First, how robust are LLM-generated dispatch functions to prompt variation? Small changes in wording, problem framing, or the amount of context provided may substantially shift the strategy distribution; a systematic prompt sensitivity study would clarify how much the current results depend on the specific templates used. Second, how should synthetic benchmarks for load balancing be constructed to capture the diversity of real-world scenarios? The two setups studied here (heartbeat staleness and noisy SRPT) are representative but not exhaustive; a richer template library spanning different arrival models, service distributions, and information structures would enable more general evaluation. Third, and most directly motivated by the RL generalization hypothesis above: can reinforcement learning with a simulator as the reward signal improve LLM load-balancing capabilities? Using the discrete-event simulation itself as a judge—scoring generated dispatch functions by the queue standard deviation or mean response time they achieve—would provide a task-specific reward signal absent from standard training, and could close the gap between occasional correct solutions and consistently good ones.

6 Conclusion

We evaluated six LLMs on two load-balancing synthesis tasks designed to reflect a common practical situation: a practitioner knows a standard algorithm exists but suspects it is underperforming, and asks an LLM for a better alternative.

The answer, across both experiments, is that the LLM should not be trusted blindly. No model reliably beats the best available baseline in median performance, and the typical LLM-generated function performs worse than the baseline it is meant to improve. Scale determines whether a model can produce valid code at all (roughly 8B parameters is the threshold) but does not predict median quality above that threshold. Larger models do occasionally discover correct solutions, but they do so inconsistently—the right answer appears in a small minority of trials, not as the expected output.

Providing an explicit description of the failure mode substantially increases the rate of correct solutions, suggesting that LLMs have the underlying capability to fix described problems but not to diagnose them from scratch. The most direct next steps are to measure how much of this gap is driven by prompt framing versus task difficulty, and to explore whether task-specific reinforcement learning—using the simulator itself as the reward signal—can close the gap between occasional correct solutions and consistently reliable ones.

References

- A. Cheng, S. Liu, M. Pan, Z. Li, S. Agarwal, M. Cemri, B. Wang, A. Krentsel, T. Xia, J. Park, S. Yang, J. Chen, L. Agrawal, A. Naren, S. Li, R. Ma, A. Desai, J. Xing, K. Sen, M. Zaharia, and I. Stoica. Let the barbarians in: How AI can accelerate systems performance research, 2025a.
- A. Cheng, S. Liu, M. Pan, Z. Li, B. Wang, A. Krentsel, T. Xia, M. Cemri, J. Park, S. Yang, J. Chen, L. Agrawal, A. Desai, J. Xing, K. Sen, M. Zaharia, and I. Stoica. Barbarians at the gate: How AI is upending systems research, 2025b.
- M. E. Crovella and A. Bestavros. Self-similarity in world wide web traffic: Evidence and possible causes. *IEEE/ACM Transactions on Networking*, 5(6):835–846, 1997. doi: 10.1109/90.650143.
- M. Dahlin. Interpreting stale load information. *IEEE Transactions on Parallel and Distributed Systems*, 11(10):1033–1047, 2000. doi: 10.1109/71.888643.
- A. Fawzi, M. Balog, A. Huang, T. Hubert, B. Romera-Paredes, M. Barekatin, A. Novikov, F. J. R. Ruiz, J. Schrittwieser, G. Swirszcz, et al. Mathematical discoveries from program search with large language models. *Nature*, 626:468–475, 2024. doi: 10.1038/s41586-023-06924-6.
- H.-C. Lin and C. S. Raghavendra. An approximate analysis of the join the shortest queue (JSQ) policy. *IEEE Transactions on Parallel and Distributed Systems*, 7, 1996. doi: 10.1109/71.491583.
- Y. J. Ma, W. Liang, G. Wang, D.-A. Huang, O. Bastani, D. Jayaraman, Y. Zhu, L. Fan, and A. Anandkumar. Eureka: Human-level reward design via coding large language models. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024.
- M. Mitzenmacher. How useful is old information? *IEEE Transactions on Parallel and Distributed Systems*, 11(1):6–17, 2000. doi: 10.1109/71.824714.
- M. Mitzenmacher. The power of two choices in randomized load balancing. *IEEE Transactions on Parallel and Distributed Systems*, 12(10):1094–1104, 2001. doi: 10.1109/71.963420.
- A. Novikov, M. Balog, T. Hubert, T. Bauber, B. Romera-Paredes, A. Fawzi, et al. AlphaEvolve: A coding agent for scientific and algorithmic discovery, 2025.
- K. Ousterhout, P. Wendell, M. Zaharia, and I. Stoica. Sparrow: Distributed, low latency scheduling. In *Proceedings of the ACM Symposium on Operating Systems Principles (SOSP)*, pages 69–84, 2013. doi: 10.1145/2517349.2522716.
- L. Schrage. A proof of the optimality of the shortest remaining processing time discipline. *Operations Research*, 16(3):687–690, 1968. doi: 10.1287/opre.16.3.687.